

Worked for 27s

可以。我刚刚按你这份 **8页新版 PPT 的实际内容**重新看了一遍。考虑到你下周是给比较大的领导 present，我建议整体控制在 **7–9 分钟**，不要把它讲成 machine learning technical review，而要让他们最后记住三件事：

We trained a model using our own historical data.

The model can already generate Low / Medium / High risk predictions.

The next step is to turn the local POC into a controlled, explainable AML decision-support prototype.

下面这版我特意把英文写得比较自然、句子不复杂，方便你直接练习，不需要背非常书面的英文。

Slide 1 — Building an AI-Driven AML Risk Scoring POC

建议时间：40–50 秒

你可以这样讲：

Good morning everyone. Today I'd like to give you a quick update on our AI-driven AML risk scoring proof of concept.

The main goal of this POC is to evaluate whether we can use our own historical client data to train a machine learning model that predicts AML risk into three categories — Low, Medium, and High.

I won't spend much time today on the technical details of model training. Instead, I'd like to focus on what we have built, how the model performs, what the current demo looks like, and where we can potentially take this next.

这一页的核心

不要一开始说：

I used Python, Pandas, XGBoost...

领导还没进入状态。

你是先告诉他们：

“我今天不是来教 machine learning，我是来展示这个 idea 已经做到什么程度了。”

转场

So first, let me quickly summarize where we are today.

Slide 2 — Executive Takeaway

建议时间：50–60 秒

这一页其实是整个 PPT 的 **Executive Summary**。

你可以说：

At a high level, we have already completed the core model development.

We took our internal historical data, prepared it for machine learning, and trained a model that can return a Low, Medium, or High risk prediction, together with the probability for each category.

So the key question is no longer only, "Can we train a model?"

We have already demonstrated that technically.

The next question is, "How can we package this capability in a way that is useful, explainable, and easy to review from an AML perspective?"

And importantly, we're positioning this as **decision support for analysts**, not as a replacement for final AML judgment.

最后这句话建议稍微慢一点：

Decision support — not automated decision making.

这句话领导大概率会比较喜欢。

转场

Let me show you what the end-to-end process looks like.

Slide 3 — End-to-End AML Risk Scoring Flow

建议时间：50–60 秒

这一页不要一个 box 一个 box 很机械地念。

你可以这样讲：

The overall concept is actually very straightforward.

We start with historical client data, including relevant KYC and AML attributes.

We prepare those fields into a model-ready format, and because the historical records already have known risk outcomes — Low, Medium, or High — the model can learn the patterns associated with those outcomes.

Once the model is trained, we can give it a new client record and ask it to generate a risk prediction and a confidence score.

So in simple terms, we're using historical decisions to identify patterns that can help support future risk assessments.

这里我建议你重点讲最后一句：

We're using historical decisions to identify patterns that can support future risk assessments.

比解释 supervised learning 好很多。

转场

And those predictions are not based on just one field or one indicator.

Slide 4 — Model Inputs

建议时间：45–55 秒

你可以说：

One important point is that the model is not looking at one single risk indicator.

It can evaluate multiple dimensions together.

For example, we can look at customer profile information, KYC and AML factors, geographic risk, and relationship or activity-related information.

The actual model uses the fields available in our dataset, but I've grouped them here into business-friendly categories rather than showing every individual column.

The value of the machine learning approach is that it can evaluate combinations of these factors and learn how they historically relate to different risk outcomes.

最后一句很重要：

It evaluates the combination of risk factors rather than treating every factor independently.

这句话体现 ML 的价值。

⚠ 你正式 present 前

Slide 4 现在还有：

Replace these examples with your actual field categories before presenting.

这一句一定要删掉。

这是 template instruction，绝对不要留给领导看到。

而且里面的 field 如果不是你真实数据里的，也建议换成你的真实 category。

转场

Of course, the most important question is whether the model actually performs well on data it hasn't seen before.

Slide 5 — Model Validation

建议时间：1–1.5 分钟

这是你整个 presentation 最重要的两页之一。

等你把自己的 Accuracy / Recall / Confusion Matrix 放进去以后，可以这样讲。

假设你的实际数据是：

- Accuracy 86%
- High Risk Recall 91%
- Medium 78%
- Low 88%

你不要机械念：

Accuracy is 86%, high recall is 91%...

而是：

We also validated the model using data that was held out from training.

Overall, the model achieved approximately **[XX%] accuracy**.

But for an AML use case, I think it's more important to look beyond overall accuracy, especially at how the model performs on the High Risk population.

In this test, the model correctly identified approximately **[XX%] of the actual High Risk cases**.

The heatmap on the right shows the actual historical risk rating versus the model prediction.

Ideally, we want most of the observations concentrated along the diagonal, which means the model prediction matched the actual historical outcome.

We can also use this view to understand where the model is making mistakes — for example, whether Medium Risk cases are sometimes classified as Low or High.

一个非常重要的表达

不要说：

The model is 90% correct.

最好说：

On the held-out test dataset, the model achieved...

因为这更严谨。

如果领导问：“Why High Risk Recall?”

你可以回答：

For AML, missing a potentially High Risk client may be more significant than incorrectly escalating some lower-risk cases, so High Risk recall is one metric I would pay particular attention to.

非常好用。

转场

Once we validated the model, the next step was to test whether we could actually use the trained model to score a client record.

Slide 6 — Current Demo

建议时间：1 分钟

这个就是你之前担心的：

目前只有 VS Code / Terminal，可不可以给领导看？

可以。

而且你的说法很重要。

你可以直接讲：

At the current stage, this is still a locally running technical POC.

I trained the model in VS Code, saved the trained model, and then created a separate prediction process that loads the model and scores a client record.

This screenshot shows an example of that process running locally.

For this sample client, the model returns probabilities across Low, Medium, and High risk and then assigns the final predicted risk category.

So although the interface today is still a terminal, the core end-to-end scoring capability is already working.

然后指向 screenshot：

For example, here you can see the model assigns **[High Risk]** with approximately **[XX%] probability**.

接下来这句话很关键：

The terminal itself is not intended to be the final user experience.

The next step would simply be to expose the same model capability through an API and eventually a lightweight user interface.

这就把：

“为什么还是 Terminal？”

提前回答掉了。

而且显得你的 roadmap 是有设计的。

不要说

Sorry, I haven't built the UI yet.

绝对不要。

换成：

At this stage, the focus was validating the model capability before investing in the interface layer.

非常专业。

转场

However, in AML, getting a prediction is only part of the problem. We also need to understand why the model made that prediction.

Slide 7 — Explainability

建议时间：1 分钟

这是第二个非常重要的页面。

你可以说：

For an AML use case, we don't want a model that simply says, "This client is High Risk," without any explanation.

We also need to understand which factors are contributing to that prediction.

That's where model explainability comes in.

Using tools such as SHAP, we can identify which features are pushing an individual client toward a higher or lower predicted risk.

For example, in this case, factors such as **[your real factors]** may be contributing more strongly to the final prediction.

This gives us a much more useful output than just a risk score, because an analyst can see both the prediction and the key drivers behind it.

然后收一句：

So the target output is not just "High Risk." It is "High Risk — and here is why."

这一句话我特别建议你背下来。

很容易让领导记住。

如果 SHAP 还没完成

千万不要假装完成。

直接改成：

Explainability is one of the next capabilities I'm currently adding to the POC.

也完全没问题。

转场

So that brings us to where the POC is today and what I think the logical next step would be.

Slide 8 — Roadmap / Next Steps

建议时间：1 分钟

这页是你的 **closing**。

不要把 roadmap 每个点都念一遍。

你可以说：

So to summarize where we are today:

Data preparation is completed.

Model training is completed.

The model can already generate predictions locally.

And we're now validating the results and preparing the explainability layer.

From here, I see the next step as packaging the existing model into a controlled prototype rather than immediately trying to build a production solution.

That could include a simple prediction API, a lightweight demo interface, and a small validation package using anonymized client cases.

This would allow AML stakeholders to review not only the model performance, but also whether the output is useful from a business and analyst perspective.

然后最后收尾：

So the main question I'd like to leave with the group is whether this is a direction we would like to continue exploring as a controlled AML decision-support POC.

说完。

停。

不要继续解释。

等领导问问题。

你整个 Presentation 串起来其实就是这个故事

你不要把它记成八页。

你脑子里面只需要记 六句话：

1. We wanted to see whether our own historical data could be used for AML risk prediction.
2. We trained a model that predicts Low, Medium, and High risk.
3. The model learns across multiple client and risk dimensions.
4. We validated the model against historical outcomes.
5. The model can already score a client locally, and we can explain why it made the prediction.
6. The next step is packaging it into a controlled AML decision-support prototype.

只要这六句话记得住，即使 presentation 中间突然忘词，你也不会乱。

大领导可能问你的问题

我特别建议你提前练这几个。

1. “Is this intended to replace the current AML risk rating process?”

你可以回答：

No. At this stage, I'm positioning it as a decision-support capability rather than a replacement for the existing process. The goal is to evaluate whether the model can

provide an additional data-driven signal to support analyst review.

2. “How accurate is it?”

不要只回答数字。

On the held-out test dataset, the model achieved approximately **XX% overall accuracy**. I'm also looking specifically at class-level performance, particularly High Risk recall, because overall accuracy alone doesn't tell us enough for an AML use case.

3. “Where did the training data come from?”

The POC uses our historical internal dataset, where the client attributes and the historical risk outcomes are already available. For the presentation and demo, any client-level information would be anonymized.

4. “How do we know why it rated someone High Risk?”

That's exactly why explainability is important. The goal is to provide the prediction together with the key factors that drove the model toward that risk category, rather than presenting the model as a black box.

5. “Is this production-ready?”

这个问题一定很可能出现。

不要说 yes。

No, I would still consider this a POC. The core model and local prediction capability are working, but before production we would need additional validation, governance, integration, monitoring, security review, and business ownership.

这个回答非常稳。

6. “What do you need next?”

你可以回答：

My recommendation would be a relatively small next step: package the existing model behind a simple API or lightweight interface, prepare several anonymized test cases, and let AML stakeholders evaluate both the prediction and the explainability output before we consider anything larger.

你 present 时有几个细节我很建议注意

第一，不要说“AI decided this client is High Risk.”

说：

The model predicted High Risk.

或者：

The model assigned a High Risk prediction.

“decided” 在 AML 场景太重。

第二，不要一直说：

XGBoost、OneHotEncoder、Pipeline、SHAP package、Python...

除非别人问。

你的主 deck 只需要出现：

Machine Learning Model

他们如果真的感兴趣，你再回答：

The current implementation uses XGBoost.

第三，你不要表现成在推一个已经完成的 production solution。

你的定位应该始终是：

We have validated enough of the technical concept that it is worth considering the next POC stage.

这会安全很多，也比较符合领导层 presentation。

我建议你练习到这个长度

Slide	时间
1 Opening	0:45
2 Executive takeaway	0:50
3 Workflow	0:55
4 Inputs	0:45
5 Validation	1:15
6 Demo	1:05
7 Explainability	1:00
8 Next steps	1:00
Total	约 7.5 分钟

如果给你 15 分钟的 meeting，也不要主动讲 15 分钟。

7–8 分钟讲完，剩下时间让领导问问题，反而更像 executive presentation。

另外我看你这份 8 页 PPT 后，有几个地方在正式给大领导 present 前我会建议再改一下，尤其是 **Slide 2** 现在有一些像“给 presenter 的 instruction”而不像最终领导版内容，**Slide 4/5/6** 也还有 **template wording**。如果你愿意，我下一步可以直接把这份 8 页 PPT 再帮你做一次 **真正的 final leadership version**：不增加页数，只把这些 placeholder/instruction 文案清掉，并把标题和 wording 全部调整成可以直接拿去 present 的版本。